

Parallelism Day

Avi Kulkarni

Dartmouth College

July 2, 2022

Goals

Understand how to use parallelism within magma

Understand how parallelism works in magma

Agenda

This Primer

Exercises + Workshop time

Resources

```
1 git clone https://github.com/a-kulkarn/magma-parallel-  
   cookbook.git
```

Why Parallelize (in magma)?

Most loops

```
1 [Do_Something(x) : x in S]
```

are sped up considerably.

Why do this *in magma*?

- ▶ Automatic management to avoid idle CPUs
- ▶ **Magma objects** passed between subprocesses

Some **intrinsic**s (builtin magma functions) have native parallelism support.

```
1 SetNthreads (n) ;  
2 SetGPU(true) ;    // Doesn't work on MacOS.
```

Things Improved by `SetNthreads (n)` :

- ▶ Matrix multiplication
- ▶ Groebner basis calculations
- ▶ Short vector enumeration in lattices

As well as anything dependent on the above.

Warning 1: Your milage per machine may vary.

```
1 SetNthreads(N); // Where N = 1, 4
2 X := Random(MatrixRing(GF(5), 10000));
3 time P := X*X;
```

Macbook pro, M1 chip, 8 cores	$N = 1$	5.450 (s)
-------------------------------	---------	-----------

Macbook pro, M1 chip, 8 cores	$N = 4$	6.870 (s)
-------------------------------	---------	-----------

Toby, AMD Ryzen Threadripper, 48 cores	$N = 1$	5.780 (s)
--	---------	-----------

Toby, AMD Ryzen Threadripper, 48 cores	$N = 4$	2.120 (s)
--	---------	-----------

Warning 2: Threads \neq processors

```
1 SetNthreads(N); // Where N = 1, 40, 1000, 4096
2 X := Random(MatrixRing(GF(5), 10000));
3 time P := X*X;
```

Toby, AMD Ryzen Threadripper, 48 cores

$N = 1$	5.780 (s)
$N = 40$	0.820 (s)
$N = 1000$	1.380 (s)
$N = 4096$	3.740 (s)

User-Implemented parallelism in Magma

Don't use Magma's User-Implemented parallelism on a
multi-user machine

(There is a security issue)

User-Implemented parallelism in Magma

- ▶ One process opens a channel

```
1 socket := Socket(...);  
2 DistributedManager(socket, jobs); // Dangerous
```

- ▶ Other processes connect to this channel

```
1 the_results := DistributedWorker(host, port,  
    your_function);
```

- ▶ Input/Output is transmitted over the channel (serialization)

Directory: cookbook/lecture-examples/

```
1 // collatz_manager.m
2 socket := Socket(: LocalHost := "localhost",
3                 LocalPort := 10000);
4 for i in [1..10] do
5     System("magma collatz_worker.m &"); // Launch and
6     detach the workers.
7 end for;
8 DistributedManager(socket, [1..10]);
9 delete socket;
```

```
1 // collatz_worker.m
2 host := "localhost"; // Same as above.
3 port := 10000;
4
5 function collatz_info(one_arg)
6     ...
7 end function;
8
9 DistributedWorker(host, port, collatz_info);
10 quit;
```

Important takeaways

- ▶ Don't need the `SetNthreads` call.
- ▶ Manager and workers are **separate processes**
 - ▶ Source code needs to be in separate files.
 - ▶ No shared memory
 - ▶ Worker is run as a script. (Or interactively)
 - ▶ Worker needs to quit.
- ▶ Host and port information must be the same
- ▶ Only one argument can be passed.

Restrictions of Magma's parallelism:

- ▶ No shared memory between processes.
- ▶ Transmitting of objects not faithful.
- ▶ Transmitting user defined objects very restrictive.

Warning 3: Orphans

An **orphan process** is a process whose parent process has terminated.

- ▶ Easy to accidentally create when parallelizing code
- ▶ Eat resources or block ports

Example: How to “rescue” orphans

Directory: `cookbook/lecture-examples/`

```
1 $ magma dyad_institute.m // Spawns children and exits
```

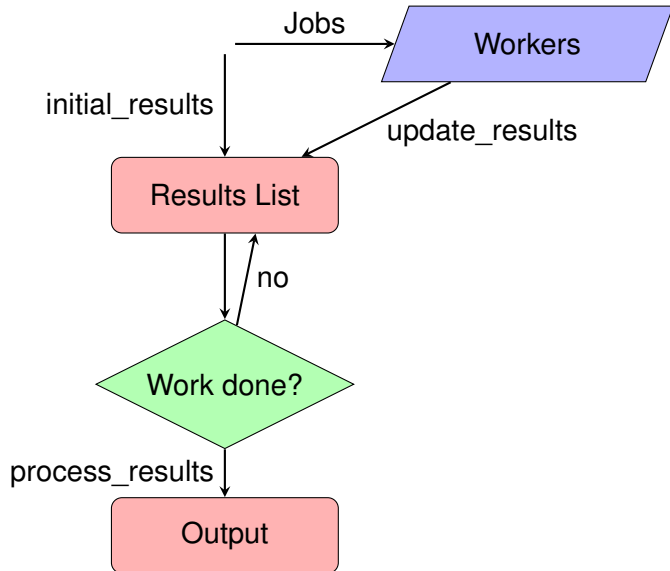
Catch-all tool:

```
1 $ top
2 $ htop // Even better if you have it.
```

Options for Distributed manager

```
1 DistributedManager(server::IO, inputs::List) -> .
2 [
3   initial_results, // Initial value for accumulator
4   update_results,  // Update on intake
5   process_results, // Post-process
6   group_tasks,     //
7   update_group,     //
8   incomplete_group //
9 ]
```

Options without Task Groups

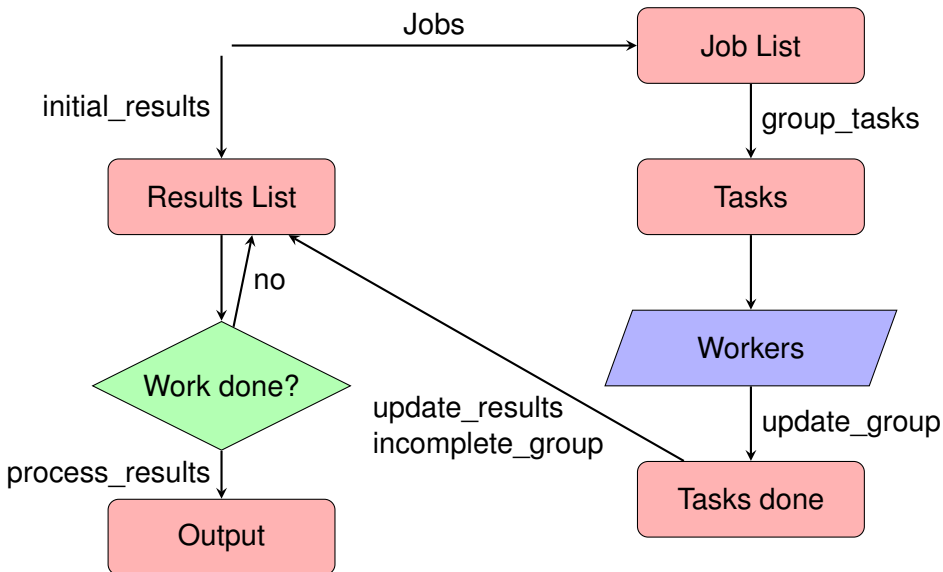


Task Groups

- ▶ Organize a Job into parallelized sub-jobs (task groups)
- ▶ Common use cases:
 - ▶ Try many things, return first-to-finish
 - ▶ Searching in a box
 - ▶ Dividing a computation among primes

```
1 DistributedManager(server::IO, inputs::List) -> .
2 [
3   initial_results, // Initial value for accumulator
4   update_results,  // Update on intake
5   process_results, // Post-process
6
7   group_tasks,      // From jobs, create task groups
8   update_group,     // Update task when workers finish
9   incomplete_group // Handle incomplete tasks.
10 ]
```

Options with Task Groups



Directory: cookbook/lecture-examples/

```
1 // squares_manager.m
2 break_into_local_tasks := function(pairs)
3     a, n := Explode(pairs); // Unpack
4     default_val := true;
5     ...
6     // Divide into tasks
7     facts := Factorization(n);
8     return default_val, [<a, p[1]^p[2]> : p in facts];
9 end function;
10
11 combine_local_info := function(item, task, tresult,
12     gresult)
13     ...
14 end function;
15 ...
16 issquare_input := [<-1, n> : n in [1..100]];
17
18 results := DistributedManager(socket, issquare_input :
19     group_tasks := break_into_local_tasks,
20     update_group := combine_local_info);
```

Summary (Thanks!)

- ▶ Parallelizing code is *sometimes* faster.
- ▶ Some Magma built-ins support parallelism.
- ▶ Distributed computing can use one or many machines.
- ▶ User-implemented parallelism is based on manager-worker model.
- ▶ The `DistributedManager` class can implement different flavours of this model.
- ▶ `top`, `htop` can get you out of trouble.
- ▶ Examples have been provided in the cookbook.

Magma also allows for a lower level interface (managing the I/O channel directly). We did not cover this today.

Warning!

The following slide is dangerous.

On toby

```
1 $ hostname // Prints "toby"
2 $ magma
3 > socket := Socket(: LocalHost:="toby",
4                     LocalPort:=10000);
5 > DistributedManager(socket, [1..10]);
```

Meanwhile... on doob

```
1 // collatz_worker.m
2 host := "toby"; // Name of host
3 port := 10000; // Note: Receiving Port needs to be open
4                //           in the firewall.
5
6 function collatz_info(one_arg)
7     ...
8 end function;
9
10 DistributedWorker(host, port, collatz_info);
11 quit;
```



```
1 $ magma collatz_worker.m // Interacts with toby
```